

# File I/O

## Advanced Systems Programming in C/C++/Rust (SoSe 21)

Harshavardhan Unnibhavi

# Topics

- What is a filesystem?
- Filesystem abstractions
- Filesystem interface
- A very simple file system
- Read/Write on the very simple file system
- Speeding up filesystem operations

# What is a filesystem?

- UNIX's defining feature: "Everything is a file"
- Network sockets, devices, data on disk...
- Filesystem is an implementation of the file interface
  - Decide how data is stored on disk
  - Performance, scalability, failures etc
  - Common operations include open, read, write
- Popular examples
  - ext4
  - procfs, sysfs
  - tmpfs

# Filesystem abstractions: File

- 10,000 foot view: Linear array of bytes
- Each file has a low-level name called the **inode**.
  - Contains metadata about file
    - Size, permissions, creation time, last access time, file type
    - Locations of data via direct or indirect pointers.
- This metadata can be accessed via the **stat** system call.
- Files are created using the **open** system call.

# Filesystem abstractions: Directory

- Each directory has an inode number.
- Contents are specific
  - Mapping between inodes and names
  - Eg: (foo, 10)
- Directory hierarchy.
  - Root directory (/)
  - "/" used to separate subsequent subdirectories and files
  - Eg: /foo/bar/hello.txt

# Filesystem interface

- Create files or directories
- Access files or directories
  - Read or write
- Delete files or directories

# Filesystem interface: Detour

- strace
  - Awesome tool
  - Refer to the previous lecture for more detail.

```
prompt> strace cat foo
...
openat(AT_FDCWD, "foo", O_RDONLY)    = 3
fstat(3, {st_mode=S_IFREG|0664, st_size=6, ...}) = 0
fadvise64(3, 0, 0, POSIX_FADV_SEQUENTIAL) = 0
mmap(NULL, 139264, PROT_READ|PROT_WRITE,
MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7fad4e7b9000
read(3, "hello\n", 131072)           = 6
write(1, "hello\n", 6hello)           = 6
read(3, "", 131072)
...
```

# Filesystem interface: Creating files

- **int open(const char \*pathname, int flags, mode\_t mode)**
  - For example to create a file use the **O\_CREAT** flag.
  - More flags and modes of operations in man pages.
- Open() returns a file descriptor (fd)
  - Allows you to perform certain operations on the file.
  - Managed by the OS on a per process basis.
    - fd is an index into a global, open file table.
    - The open file table has one entry for each open call.
  - Every process has 3 file descriptors: stdin(0), stdout(1), stderr(2).



# Filesystem interface: Reading and writing to files

- **ssize\_t read(int *fd*, void \**buf*, size\_t *count*)**
  - Return number of bytes read
- **ssize\_t write(int *fd*, const void \**buf*, size\_t *count*)**
  - Return number of bytes actually written
- **off\_t lseek(int *fd*, off\_t *offset*, int *whence*)**
  - Change read or write offset into the file
  - Tracked by the OS for each open file.
- **int fsync(int *fd*)**
  - Force write all dirty data to disk
  - Useless sometimes: <https://lwn.net/Articles/752063/>

## Filesystem interface: More calls to consider

- **int rename(const char \**oldpath*, const char \**newpath*)**
- **int stat(const char \**path*, struct stat \**buf*)**
- **int mkdir(const char \**pathname*, mode\_t *mode*)**
- **int readdir(unsigned int *fd*, struct old\_linux\_dirent \**dirp*,  
                  unsigned int *count*)**
- **int unlink(const char \**pathname*)**
- **int rmdir(const char \**pathname*)**

# Filesystem interface: Hard and soft links

- Hard links
  - `ln [TARGET] [LINK NAME]`
  - Another name for the same inode
  - Links incremented by one
  - File data deleted only if number of links is 0.
  - Cannot hard link to directory or files in other disk partitions.
- Symbolic or soft links
  - `ln -s [TARGET] [LINK NAME]`
  - Symbolic link is a file itself containing pathname of the target file as its data

# Filesystem interface: Make and mount


- **mkfs**
  - Takes as input a device and a filesystem type
  - Write an empty filesystem onto that partition
- **mount**
  - Make the filesystem available within the filesystem tree at an existing directory
  - Directory called mount point.

# A very simple file system: Assumptions

- Assumptions
  - The storage device presents itself as a set of blocks that are logically addressable from 0 to  $N-1$ , where  $N$  is the total number of blocks.
  - Each block size is 4KiB.
  - Total number of blocks is 8.
  - Inode is 256 bytes in size.

# A very simple file system: Organization

- 4 data blocks
- 1 inode block
- 1 block to store inode bitmaps
- 1 block to store data bitmaps
- 1 block to store the superblock of the filesystem
  - Contains filesystem metadata
  - Location is known
  - OS first reads this block before mounting the file system.



# A very simple file system: Inode

- Contains file metadata
  - Type of file
  - Size of file
  - Blocks
  - Protection information
  - Time information
- Data blocks information
  - Direct pointers
  - Indirect pointers

# A very simple file system: Directory

- Inode whose type is "directory".
- Inode contains pointers to data blocks
- Directory data blocks contain mapping between
  - File inode number
  - File name



# A very simple file system: Free space management

- Bitmaps are used to track free space
  - Set bit indicates that the inode or data block is used up.
- Inode bitmap tracks free inodes
- Data bitmap tracks free data blocks.

# A very simple file system: Read operation

- Open /foo/bar
- Read from the file
- Close the file
- Assume file is of 8KiB I.e 2 data blocks

# A very simple file system: Read operation

	Data bitmap	Inode bitmap	Root inode read	Foo inode read	Bar inode read	Root data read	Foo data read	Bar data[0] read	Bar data[1] read
open(bar)									
read()									
read()									

# A very simple file system: Write operation

- Create /foo/bar
- Write to the file
- Close the file
- Assume 4KiB of data, i.e 1 block, is written to the file.

# A very simple file system: Write operation

	Data bitmap	Inode bitmap	Root inode	Foo inode	Bar inode	Root data	Foo data	Bar data[0]
create( /foo/bar)		read write	read	read		read	read  write	
write()	read write			write	read  write			write

# A very simple file system: Caching

- Poor performance
  - Request same data blocks from disk repeatedly
  - Slow operation compared to reading from memory
- Solution
  - Cache blocks in main memory
  - Strategies to decide which blocks to keep and which to evict from memory
    - LRU
  - Modern OSes integrate virtual memory pages and file system pages into a unified page cache.

# A very simple file system: Caching

- Write buffering
  - Buffer writes in main memory instead of immediately persisting them.
  - Advantages:
    - Batch updates
    - Better scheduling of I/O
      - Sequential faster than random
    - Possibility of avoiding writes

See you at the Q&A