

Introduction in FUSE filesystems


Charalampos Mainas



What is FUSE


- A recap of filesystems
 - Part of the OS
 - An abstraction layer
 - Decides how data are stored and retrieved on disks
- FUSE: Filesystem in Userspace
 - Software layer in Userspace
 - Allows the creation of custom filesystems
 - Kernel code stays untouched

How does FUSE work

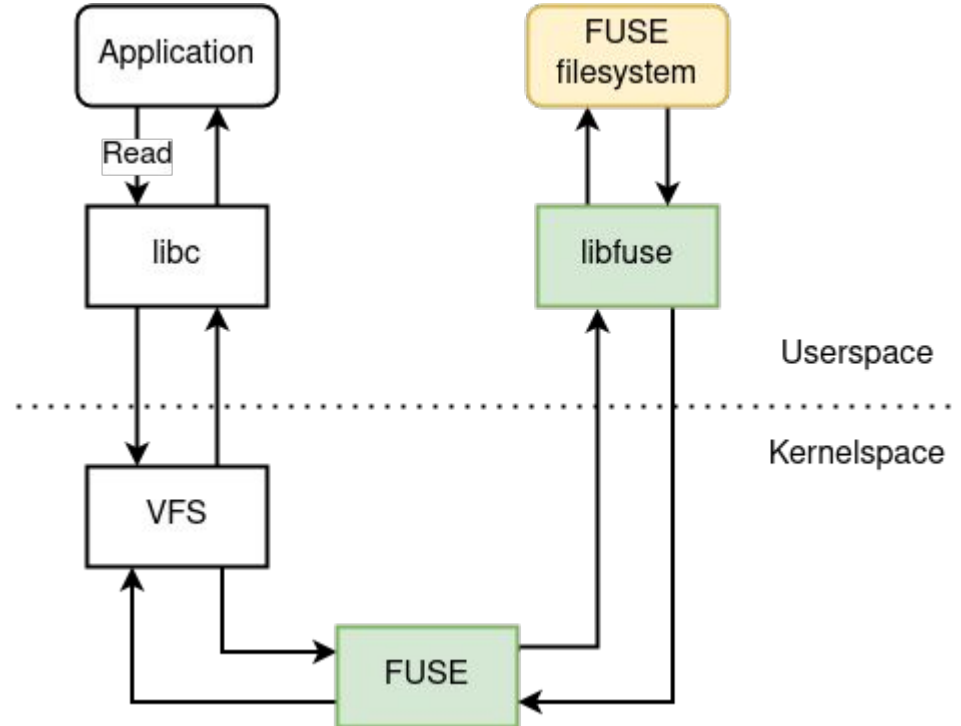


Normal filesystem

How does FUSE work



Normal filesystem



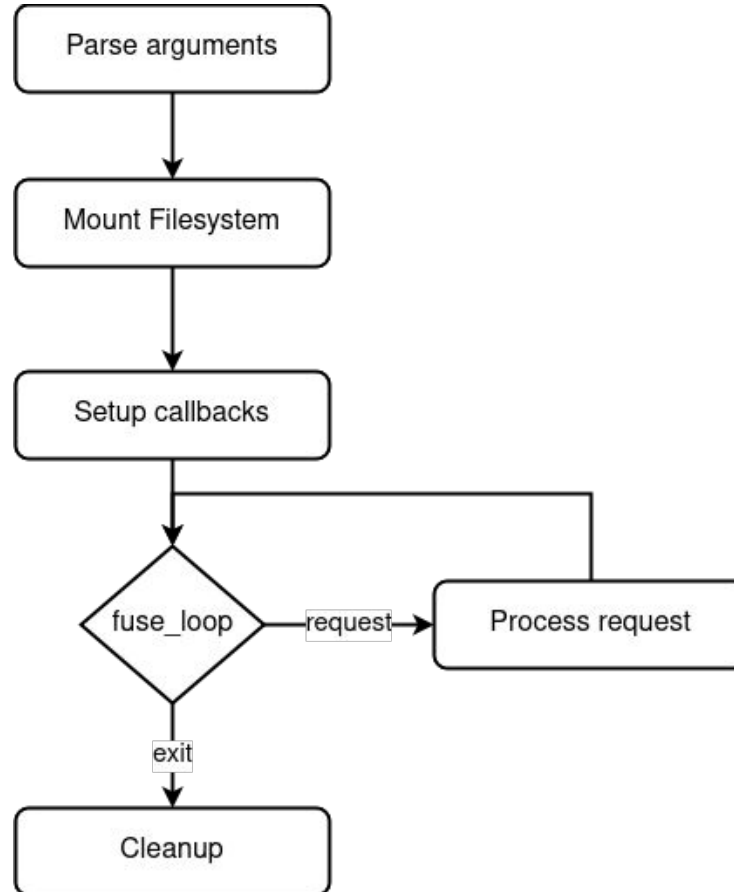
FUSE filesystem

- Why a filesystem in userspace?
 - Shorter development cycle
 - Easier development, OS-agnostic
 - Does not affect other parts/services
 - Safer usage of untrusted filesystems
- Real world use cases:
 - On-disk filesystems: NTFS, retro-fuse etc.
 - Network-based filesystems: MinFS, SSHFS, etc.
 - Layering filesystems: EncFS, FuseCompress etc.
 - Archive, backup filesystems: Atlas, Borg etc.

How to use FUSE

- A FUSE application is a typical user-space program
 - Applications define how to handle filesystem operations
 - Interaction with libfuse to register the operations
 - Libfuse invokes the application defined operation upon a request
- API from libfuse
 - Callback mechanism for binding user-defined functions with operations
 - High-level API -- path level, synchronous
 - Low-level API – inode level, asynchronous

FUSE application workflow



- FUSE allow the user to specify how a file operation will get handled
 - User implements a set of functions to handle a file operation
 - User fills the *struct fuse_operations* or *struct fuse_lowlevel_ops* with the respective function implementations
 - User passes the struct to libfuse and upon a request libfuse calls the respective function

```
static struct fuse_operations my_fuse_ops = {  
    .getattr    = my_fuse_getattr,  
    .read       = my_fuse_read,  
    .write      = my_fuse_write,  
    .readdir    = my_fuse_readdir,  
    .mkdir      = my_fuse_mkdir,  
    .mknod     = my_fuse_mknod,  
    .open       = my_fuse_open,  
    .create     = my_fuse_create,  
    .readlink   = my_fuse_readlink,  
    .symlink    = my_fuse_symlink,  
};
```


- **lookup (only in the low-level API):**
 - Searches the directory entry specified by last parameter and returns its attributes
- **getattr:**
 - Get the attributes of a file
- **read:**
 - Reads data from an open file
- **write:**
 - Writes data to an open file
- **readdir:**
 - Read a directory
- **mkdir:**
 - Create a directory

FUSE operations

- **open:**
 - Opens a file
- **mknod:**
 - Create a file node
- **create:**
 - Create and open a file
- **symlink:**
 - Create a symbolic link
- **readlink:**
 - Read the target of a symbolic link

- In case a request fails an appropriate error should get returned
- Useful error codes:
 - ENOSYS: Function not implemented
 - EPERM: Operation not permitted
 - EACCESS: Permission denied
 - ENOENT: No such file or directory
 - EIO: I/O error
 - EEXIST: File exists
 - ENOTDIR: Not a directory
 - ENOTEMPTY: Directory not empty

- Implement a in-memory filesystem using FUSE
 - Everything is stored in the process's memory
 - Choose the high or low level API of FUSE
 - Implement at least the previously mentioned operations
 - Maximum file name length of 255 ascii characters.
 - Maximum file size of 512 bytes.

- Documentation for low-level API:
https://libfuse.github.io/doxygen/fuse__lowlevel_8h.html
- Documentation for high level API:
https://libfuse.github.io/doxygen/fuse_8h.html
- FUSE can parse command line arguments
 - Check `fuse_parse_cmdline` (low level API) and `fuse_main` (High level API)
- Do not get scared with mounting the filesystem
 - `libfuse` will mount it for you
 - `fuse_mount` for low level API
 - `fuse_main` for high level API
- Use `-f` option to execute your application in foreground
 - Useful when debugging, especially with `printfs`
- Unmount a mounted FUSE filesystem with
 - `fusermount -u <mountpoint>`

Thank you for listening!
see you in the Q&A session