

A practical course on

Advanced systems programming in C/Rust

Dimitrios Stavrakakis




Today's topic!

Memory Management

- Memory organisation
 - Physical, virtual memory
- Virtual memory management
 - paging, segmentation
 - heap, stack
- Memory allocation
 - `sbrk()`, `malloc()`, `free()`, `mmap()`
 - selection & allocation policies
 - fragmentation

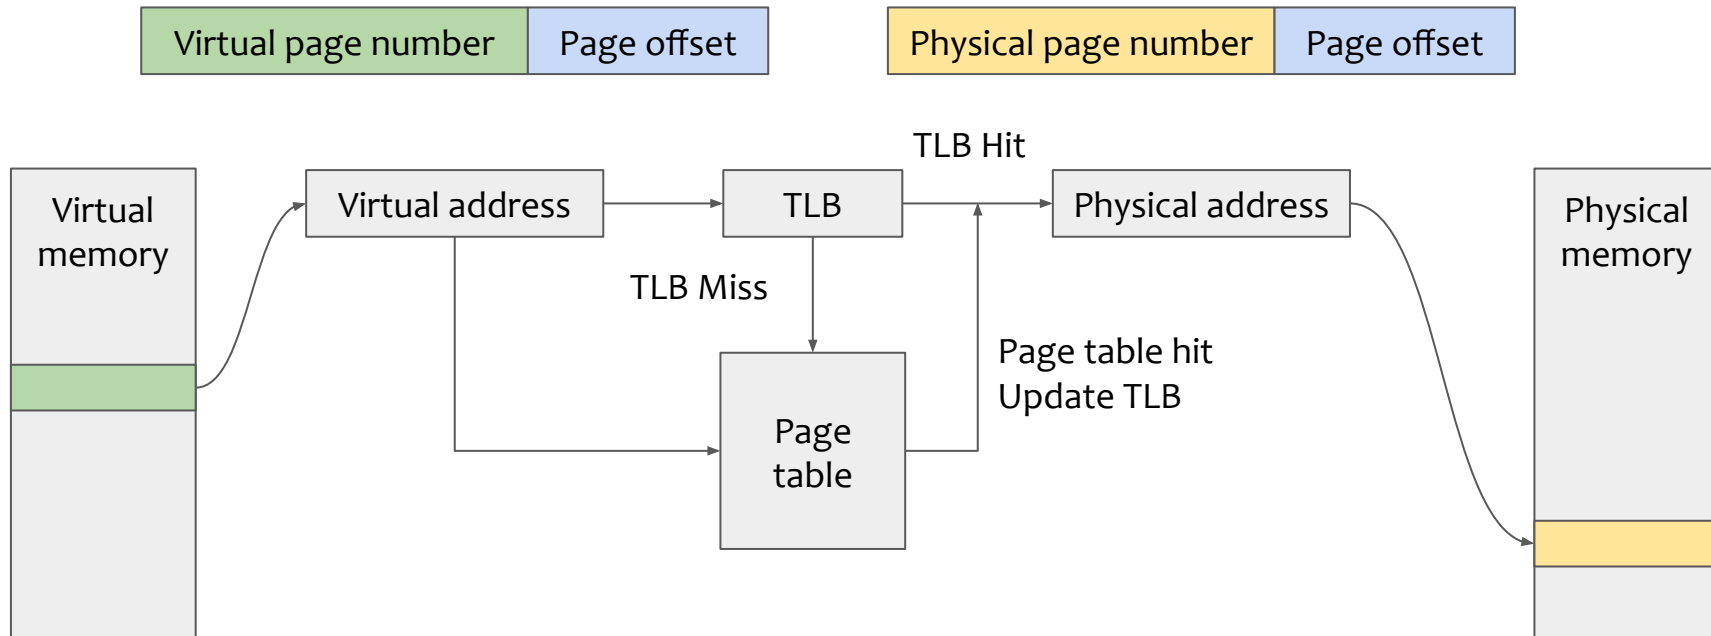
Physical & Virtual memory

- Physical memory:
 - The actual DRAM device
 - Byte-addressable directly accessed
 - Limited in size
- Virtual memory:
 - Memory management technique
 - Convenience to the programmers
 - Managed by the OS
 - Address translation required



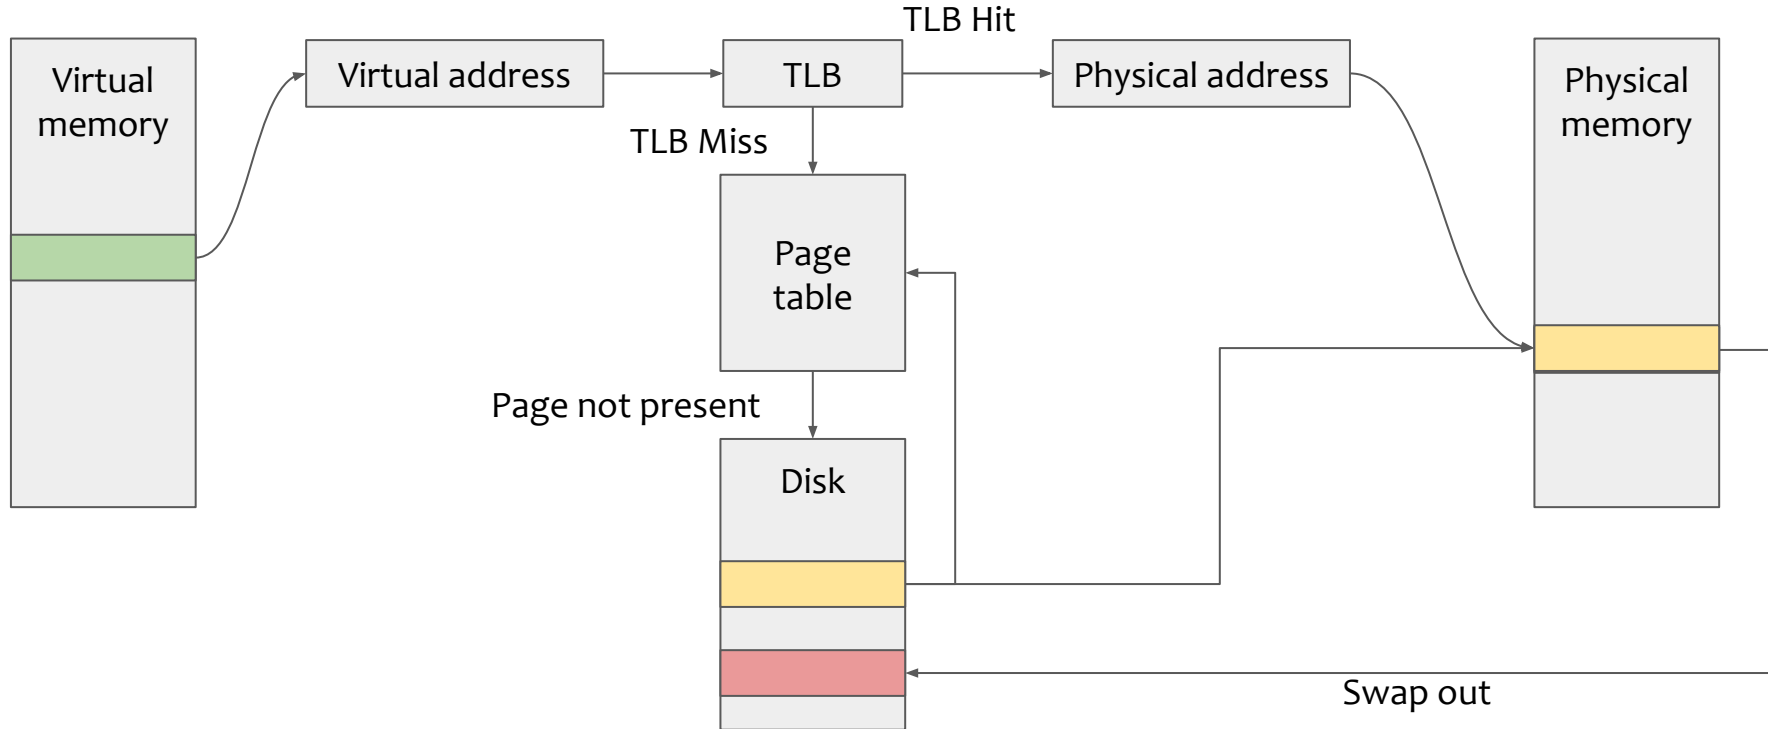
- Virtual address space is not mapped 1 by 1 to physical memory
- Need for address translation
- 2 most common approaches:
 - Paging
 - Segmentation

- Virtual address space of a process is mapped to the physical memory in pages
 - Typical page size is 4KB
- How address translation lookup process works:



Page faults

- What happens when the requested page is not found during the translation?
 - Page fault! - potentially expensive I/O operation



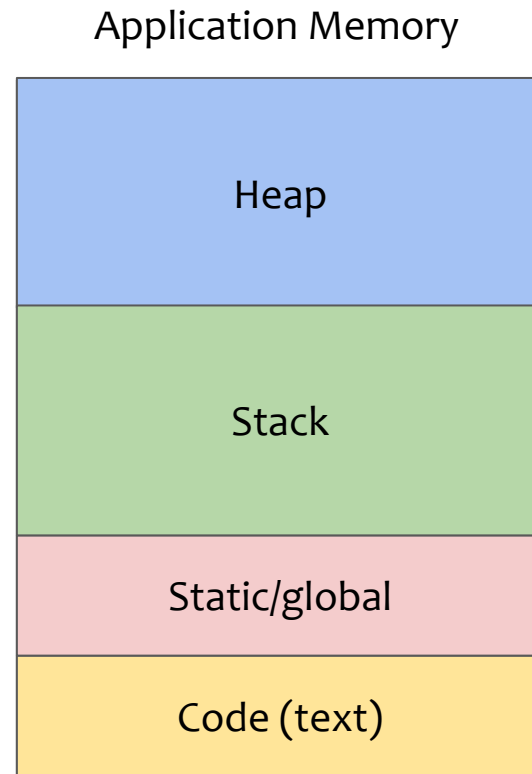
Segmentation

- Virtual address space is divided into segments of variable length
 - Each segment represents different logical address space of the program
- Segments of a process are loaded into memory at run time - not necessarily contiguously
- Similar translation with paging
 - Segment table instead of page table

Segment ID	Segment size	Physical memory address
0	100	400
1	200	800
2	100	600

Stack & Heap

- Stack
 - Contiguous block of memory
 - Allocation & deallocation performed by the compiler
 - Fixed size during run-time
- Heap
 - Heap memory can be allocated in any random order
 - Allocation & deallocation is a programmer's task
 - Heap can grow during the lifetime of an app



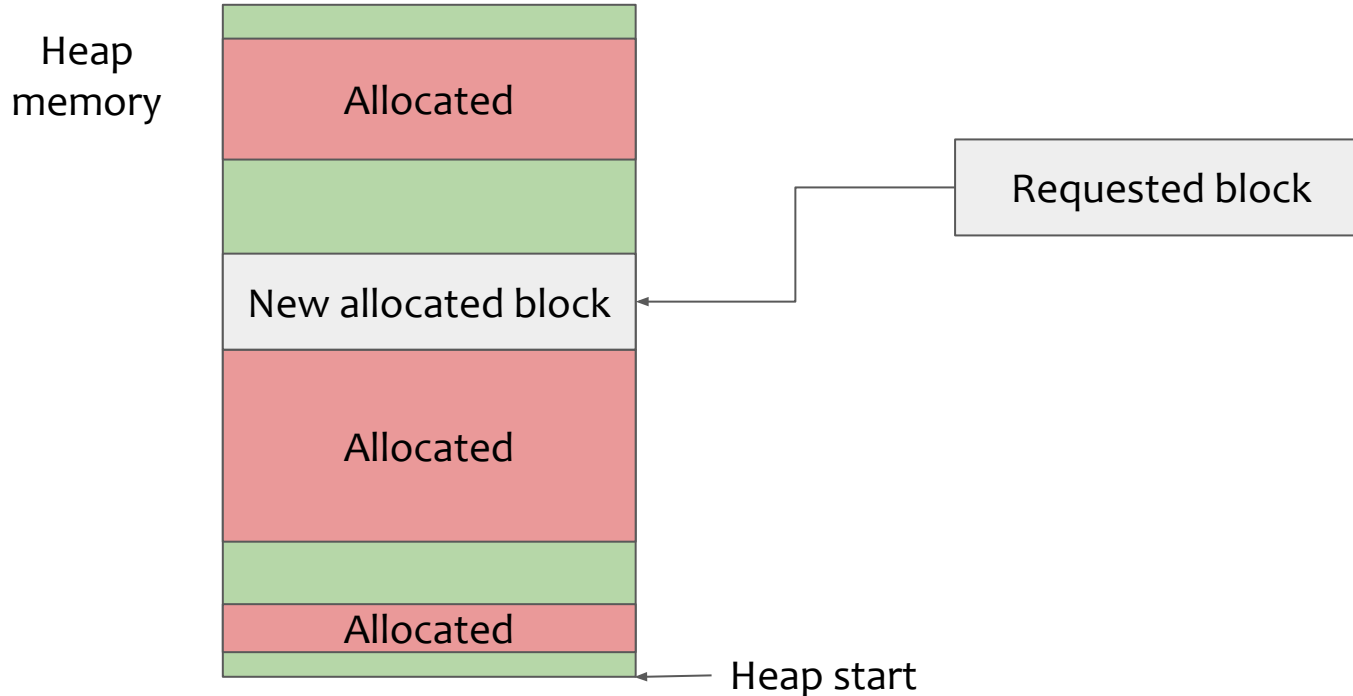
Heap memory allocation interface

- **sbrk(intptr_t increment) :**
 - Extends the heap region by the requested increment size
- **malloc(size_t size) :**
 - Allocates a memory block of at least the requested size
- **realloc(void *ptr, size_t size)**
 - Reuse the same block if size fits otherwise allocate a new block and copy the data
- **free(void *ptr) :**
 - Mark the occupied block pointer by ptr as free & reusable
- **mmap() - munmap()**
 - Maps files or devices directly into virtual memory address space
 - Uses demand paging
 - Byte-to-byte correlation

Selection policies

- **First fit:**

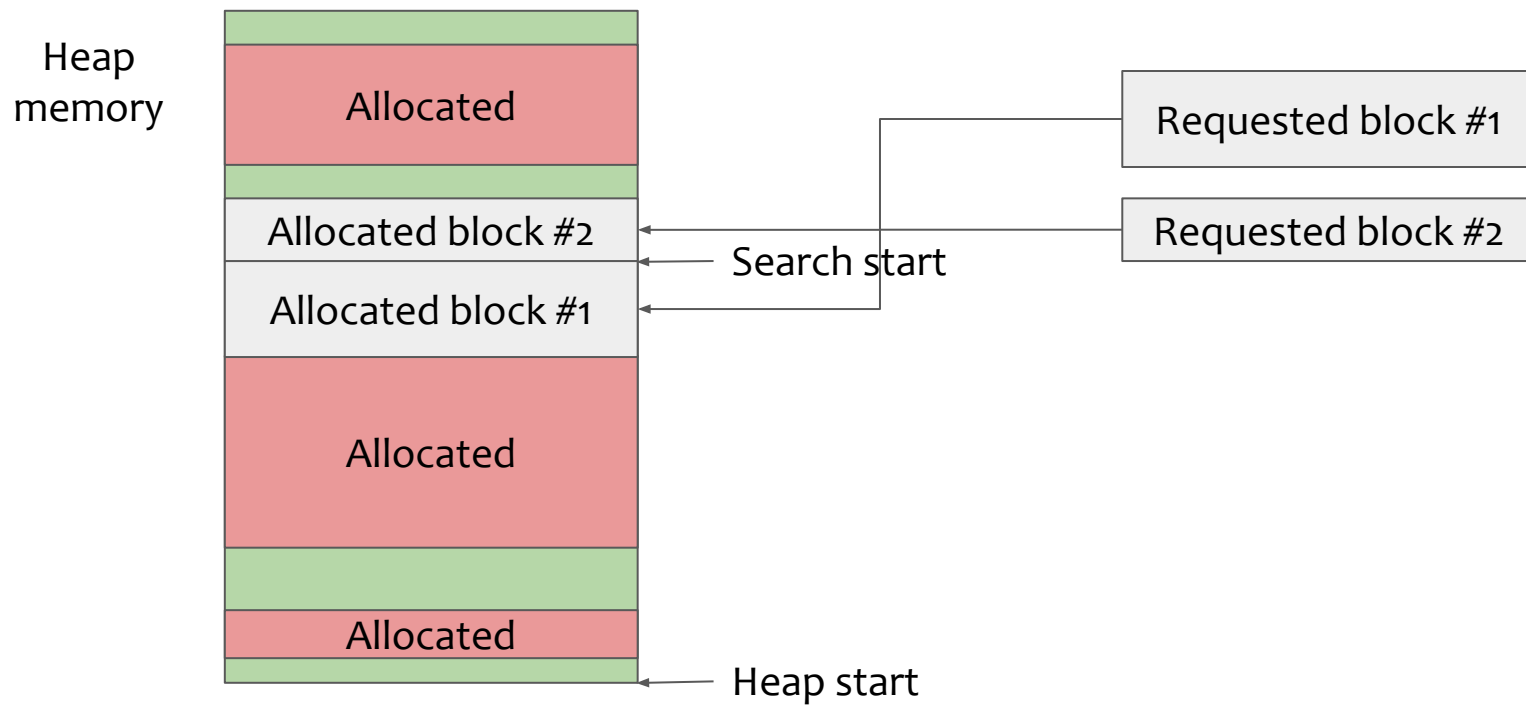
- Commence the search from the start address of the heap
- Allocate the first block that can accommodate the requested memory size



Selection policies

- **Next fit:**

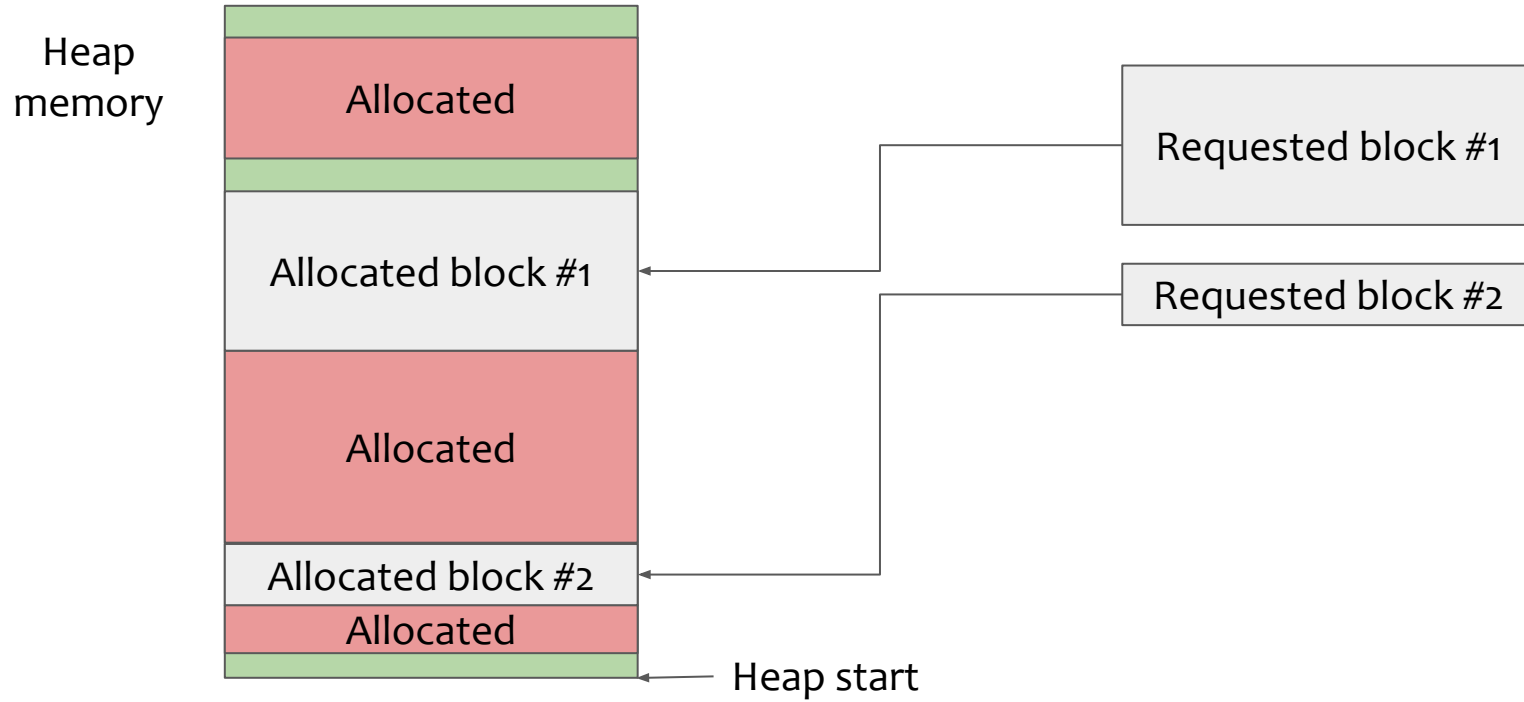
- Commence the search from the address of the most recently allocated block
- Allocate the first block that can accommodate the requested memory size



Selection policies

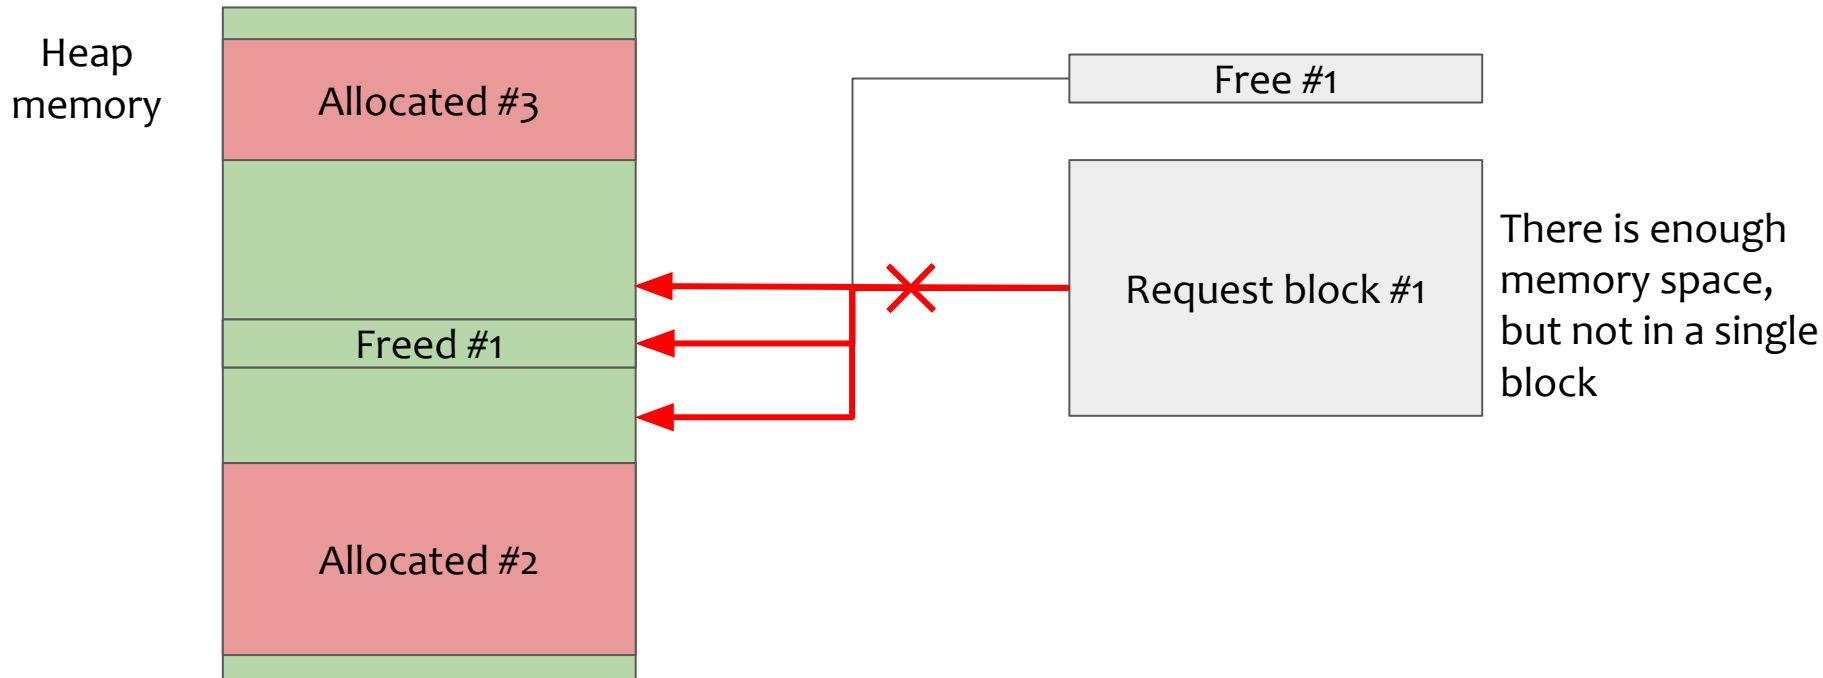
- **Best fit:**

- Commence the search from the start address of the heap
- Allocate the smallest free block that can accommodate the requested memory size



Fragmentation

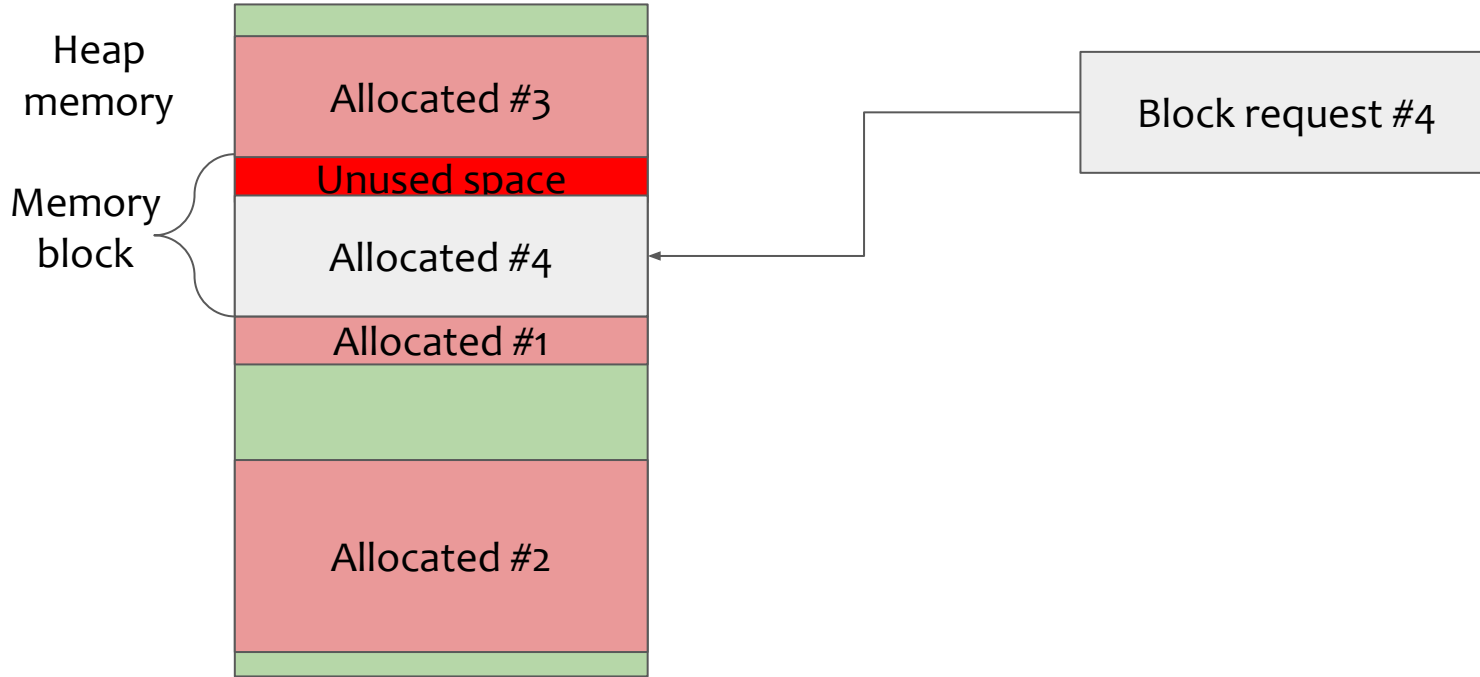
External fragmentation problem



Coalescing: merging adjacent free blocks

Fragmentation

Internal fragmentation problem



Internal fragmentation: depends on the allocation strategy

Possible allocation strategies

- Free Bitmaps
 - The state of a memory block is represented by a bit in a bitmap
- Free lists
 - Linked list that maintains the location and size of the free memory blocks
- Segregated fits
 - Maintain separate free lists of size classes referring to different sized blocks
- Buddy system allocation
 - Each block may be split into two sub-blocks (buddies) being power of 2 in size

Tasks:

- Implement your own heap allocation functions
 - `malloc()`, `realloc()`, `free()`, `calloc()`
- Make use of different allocation policies to lessen fragmentation and efficiently use the allocated heap size
 - `first_fit`, `next_fit`, `best_fit`
- Extend your allocator to make it suitable for multi-threading environment

Useful tools:

- Valgrind, AddressSanitizer, ThreadSanitizer

Thank you for listening!

See you in the Q&A session