

A practical course on

Advanced systems programming in C/Rust

Dr. Atsushi Koshiba
Marcel Schneider



Today's topic!

Performance

- Performance
 - What is the performance?
- Optimization
 - Cache locality
 - Multi-threading
 - SIMD instructions
- Measurement
 - `clock_gettime()`
 - `perf`
- Assignment

Metrics that quantify the effectiveness/efficiency of computer systems

- **Processing time (s)**
 - Simulation, numerical computation
- **Throughput (Op/s, B/s, Tx/s)**
 - HPC, database, network, storage
- **Latency (s, cycles)**
 - Database, network, storage, HFT
- **Memory usage (Bytes, pages)**
 - Embedded systems
- **Energy usage (V, A, W, W/op)**
 - Embedded systems, laptops, smartphones, tablets
- **Binary size (Bytes)**
 - Embedded systems, loading from remote
- **Compiling time (s)**
 - Testing, development

How to improve the performance

Many features affecting the performance of modern computer systems

- Developers need to carefully (explicitly) design their code to utilize the features

Hardware:

- Pipelining
- **CPU Caches**
 - Cache size
 - Caching protocol
 - False sharing
 - Cache misses
 - Prefetcher
 - Page misses
- **SIMD instructions**
- Branch predictor
- Interrupts
- Simultaneous multithreading
- Dynamic frequency scaling
- NUMA

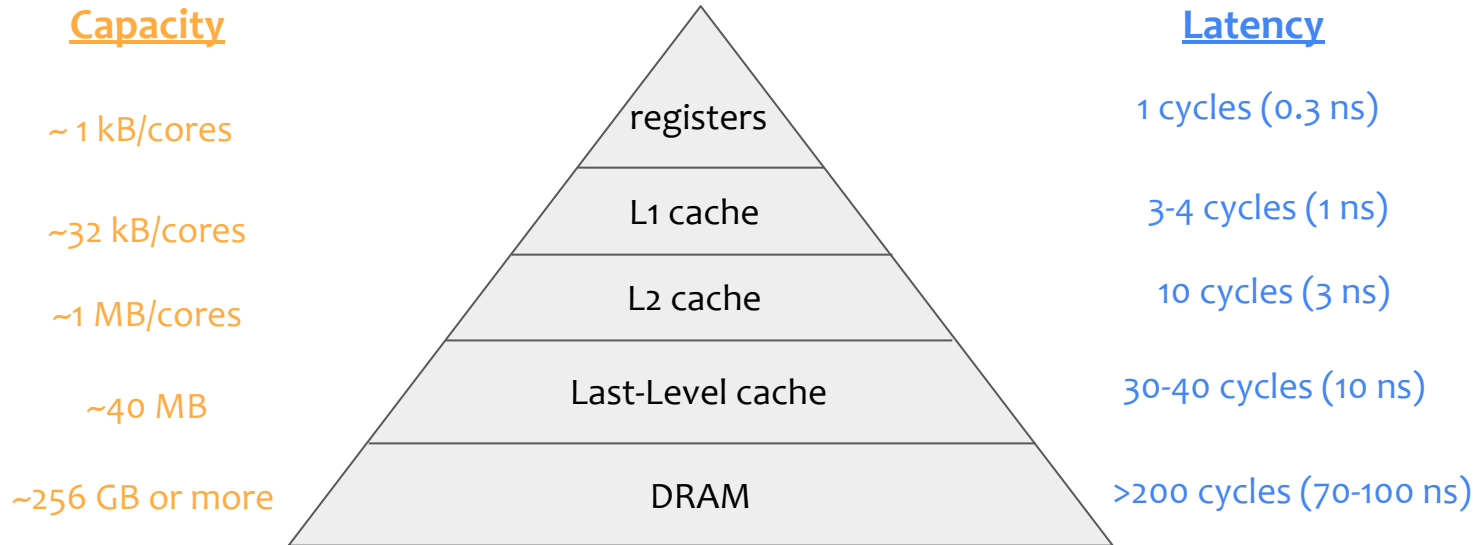
Software:

- **Multi-threading**
- Software Caches
 - Page cache
- Scheduler
 - Energy aware scheduling
 - Real time scheduling
 - Overcommitment
 - Process priority
- Context switch
- Swapping
- Compiler/Linker optimizations
- ASLR
- CPU affinity
- CPU governor

1. CPU cache

DRAM accesses (~ 100 ns) are much slower than CPU instructions (< 1 ns)

- Multiple caches (L1, L2, L3/LLC) mitigate memory access latencies
 - Memory data accessed by CPUs are fetched to caches
 - Every cache miss invokes fetching a cache line of data (64 Bytes)



1. CPU cache

Cache locality

- Programs have a tendency to access the same set of data repeatedly
 - Spatial locality — location accessed recently is likely to be accessed again
 - Temporal locality — location nearby accessed recently is likely to be accessed
- The cache holds data recently accessed and neighbors of the data

Prefetchers


- Run independent from CPU cores
- Fetch data into caches in advance according to recent memory access patterns
 - effective for sequential accesses

Use case: cache locality


Naive matrix multiplication

- Matrix A, B are stored in row-major order
- Stride accesses to B elements over rows hurt cache locality
 - Cache misses occur frequently if matrix size is large enough

```
for (int i = 0; i < rows; ++i)
  for (int j = 0; j < columns; ++j)
    for (int k = 0; k < lengths; ++k)
      C[i][j] += A[i][k] * B[k][j];
```



Matrix A




Matrix B

Use case: cache locality

Make the code cache-friendly

- Matrix D is the transpose of B
- Both accesses to A and D elements are sequential


```
matrix_transpose(B, D);  
for (int i = 0; i < rows; ++i)  
    for (int j = 0; j < columns; ++j)  
        for (int k = 0; k < lengths; ++k)  
            C[i][j] += A[i][k] * D[k][j];
```



2. Multi-threading

Use multiple CPU cores to execute a single task (studied in Task 4)

- Split data into small portions
 - avoid to hurt cache locality
- Create threads to compute each portion of data concurrently



3. SIMD instructions

Special instruction sets for Single Instruction Multiple Data (SIMD)

- Processing multiple data with a single instruction
 - SIMD registers (128-, 256-, 512-bit width) are integrated into CPU cores
 - Intel: Streaming SIMD Extensions (SSE), Advanced Vector Extensions (AVX)
 - ARM: Scalable Vector Extension (SVE)

Scalar operations

$$\begin{array}{lcl} \text{op 1} & \boxed{A_1} \times \boxed{B_1} = \boxed{C_1} \\ \text{op 2} & \boxed{A_2} \times \boxed{B_2} = \boxed{C_2} \\ \text{op 3} & \boxed{A_3} \times \boxed{B_3} = \boxed{C_3} \\ \text{op 4} & \boxed{A_4} \times \boxed{B_4} = \boxed{C_4} \\ & \vdots & \end{array}$$

SIMD operations

$$\begin{array}{lcl} \text{op 1} & \left\{ \begin{array}{c} \boxed{A_1} \\ \boxed{A_2} \\ \boxed{A_3} \\ \boxed{A_4} \end{array} \right\} \times \left\{ \begin{array}{c} \boxed{B_1} \\ \boxed{B_2} \\ \boxed{B_3} \\ \boxed{B_4} \end{array} \right\} = \left\{ \begin{array}{c} \boxed{C_1} \\ \boxed{C_2} \\ \boxed{C_3} \\ \boxed{C_4} \end{array} \right\} \\ & \vdots & \end{array}$$

SSE2 instruction sets

SSE2 supports 128-bit operations. See Intel Intrinsics Guide [1]

- Load/store inst.
 - `__m128 _mm_load_ps (float const* mem_addr)`
 - `void _mm_store_ps (float* mem_addr, __m128 a)`
- Arithmetic inst.
 - `__m128 _mm_mul_ps (__m128 a, __m128 b)`
 - `__m128 _mm_add_ps (__m128 a, __m128 b)`
- Conditional inst.
 - `__m128 _mm_cmpeq_ps (__m128 a, __m128 b)`
 - `int _mm_movemask_ps (__m128 a)`

Hardware Counter

- Performance Monitoring Counters (PMCs) in CPU cores
- Monitor how many times a specific event occurred in a certain period of time
 - e.g., elapsed time (cycles), retired instructions, cache misses, CPU stalls

Tools

- **Linux system calls (e.g., `clock_gettime`)**
- **perf**
- Intel VTune
- AMD μ Prof

- **int clock_gettime(clockid_t clk_id, struct timespec *spec)**
 - Get the current time of the clock in nanosecond time resolution
 - `clk_id` specifies a particular clock
 - `CLOCK_REALTIME` : system-wide real time clock. Changeable by configuration (e.g., NTP)
 - `CLOCK_MONOTONIC`: unchangeable monotonic time clock. **Recommended**
- **perf [OPTIONS] COMMAND [ARGS]**
 - Trace specific events through PMCs
 - `perf list` shows traceable events
 - `perf stat` runs a program and gather performance counter statistics
 - E.g., tracing LLC misses of all CPU cores during `ls` command:

```
$ sudo perf stat -a -e LLC-load-misses,LLC-store-misses ls
```

Optimize the execution time of two simple programs

- Matrix multiplication
- Mandelbrot set

Tips

- Make the code cache-friendly
- Parallelize computations with SIMD instructions, multi-threading
 - Up to four CPU cores are available
 - Please only use SSE/SSE2 instructions

