

A practical course on

Advanced systems programming in C/Rust

Jörg Thalheim



Today's topic!

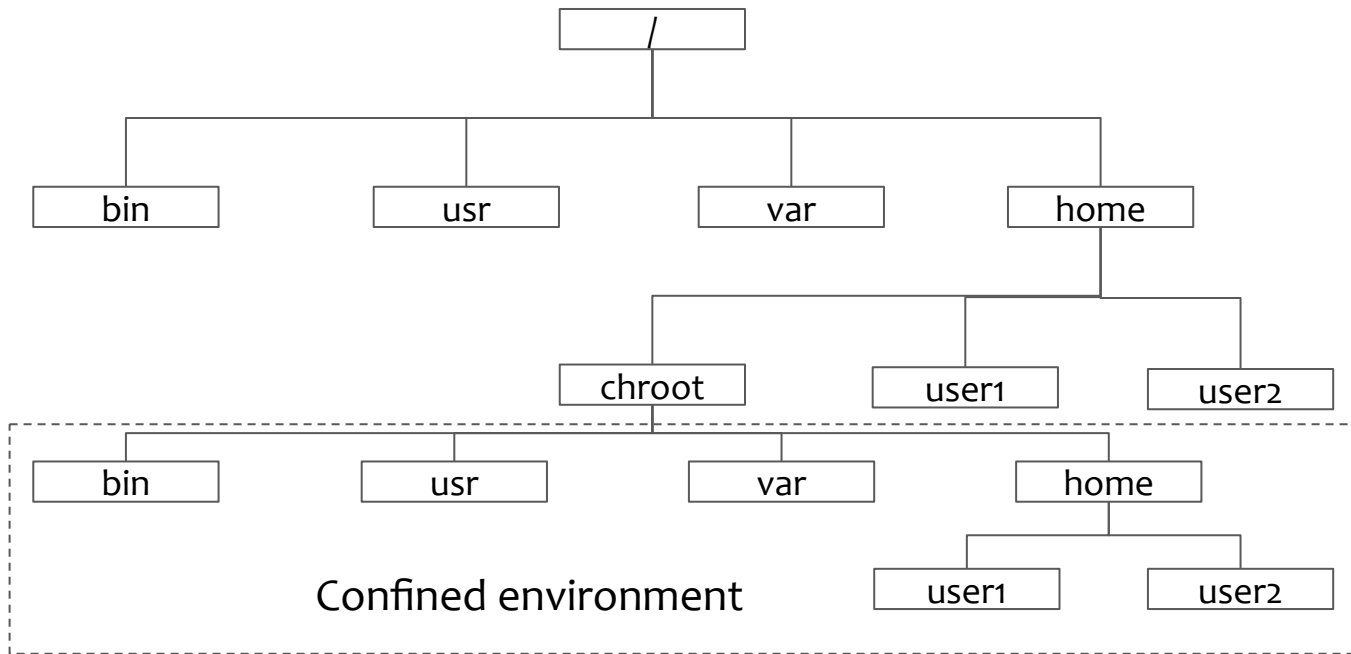
Containers

- Resource isolation only enforced in software - the operating system
 - Granularity: Process level
 - Only one OS unlike hardware virtualisation
 - Flexible resource sharing between processes
- Present in all general purpose operating system:

Unix
“chroot”

Chroot: Origins of Containers

- **Change Root** -> makes a subdirectory the new root of a process
- syscall: `chroot()`
- Useful to contain unprivileged processes (i.e. postfix mail demon, sshd)
- Root users can easily undo this



- Resource isolation only enforced in software - the operating system
 - Granularity: Process level
 - Only one OS unlike hardware virtualisation
 - Flexible resource sharing between processes
- Present in all general purpose operating system:

Unix
“chroot”


macOS
“Sandbox Apps”

Solaris
“Zones”

Windows:
“Process
Isolation”


FreeBSD
“Jails”

Our focus


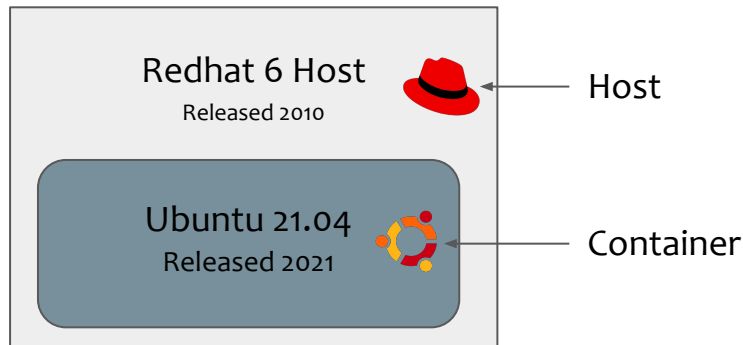


OS-Virtualisation on Linux: Container

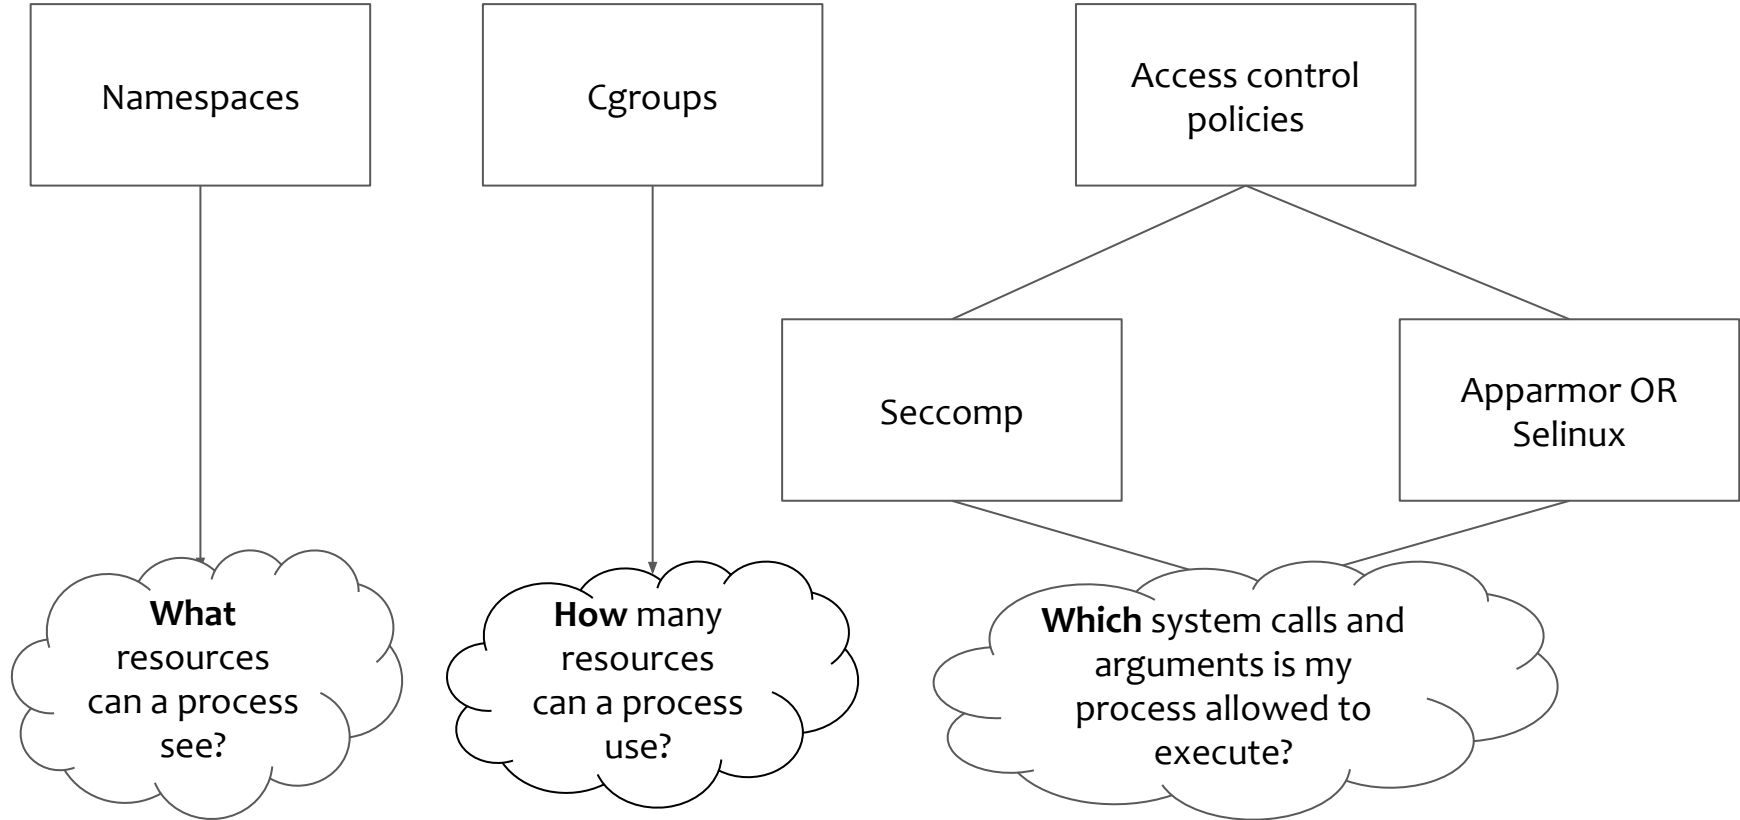
Why did it became so popular on Linux?



- (Very) stable syscall interface
 - Kernel devs: “Don't Break Userspace”
- > libc can be replaced & distributed independently (unlike i.e. FreeBSD jails)
- Convenient application packaging & distribution: Docker -> Kubernetes:
 - Less dependency hell & rollbacks



Linux APIs



- Control visibility of kernel resources (mounts, network interfaces, time etc)
- System calls and kernel subsystems needs to be aware of namespaces:
 - Stored in process context: `task_struct->nsproxy`
- Pid 1 (init process) starts with initial namespace
- Children inherit namespace from their parents

Create namespace

System call: `unshare()` or
`clone()`
Example: `$ docker run`

Switch namespace

System call: `setns()`
Example: `$ docker exec`

Inspect namespace

Interface: `/proc/self/ns`
Example: `$ ls -la /proc/self/ns`

See “man namespaces” for further information (well written)

Namespaces types

- 8 namespaces so far, but list is growing:
 - Cgroup, IPC, Network, Mount, PID, Time, User, UTS

Mount namespace

- Isolates mountpoints seen in a namespace
- Useful in combination with chroot

PID

- Isolates what processes are visible
- Allows containers to have their own Init process (Pid1)


Network namespaces

- Own routing table and set of network interfaces and firewall
- Often combined with veth-interfaces to provide virtual ethernet interfaces

User namespaces

- Allows to remap a user / group ids to different ranges
- Important for unprivileged containers (rootless), i.e.
 - Uid 1000 on the host becomes uid 0 (root uid) in the namespace

- Hierarchical resource groups:
 - Child inherits cgroup from parent
 - 9 different subsystems for different purposes
- Limits how much resources processes
 - Cpu schedule slices
 - Memory limits
 - Block I/O weight
 - Freeze processes (freezer)



- Also reliable process tracking:
 - SystemD/Docker puts each service/container in a cgroup
 - Member's of cgroup == all childs belonging to service/container
- Major API change: Cgroup v1 -> Cgroup V2
 - unified cgroup instead of separated by subsystem
- Interface through a filesystem: /sys/fs/cgroup/ (in this example cgroup v2)

Create cgroup
\$ mkdir
/sys/fs/cgroup/foo

Join a cgroup
\$ echo \$PID >
/sys/fs/cgroup/foo/cgroup
.procs


**See current cgroup
membership**
\$ cat /proc/self/cgroup

- **Goal:** restrict (root) user further by restricting system calls it can execute
- Capabilities
 - More granular permissions instead of privileged vs unprivileged processes
 - Examples:
 - CAP_CHOWN: change uid/gid of arbitrary files
 - CAP_NET_ADMIN: network configuration
- Seccomp/seccomp-ebpf
 - Filter programs that prevent system calls based on arguments
 - Example: disallow chown() with setuid bit set.
- MAC:
 - Application firewall: More complex policies/profiles to allow/disallow what files/resources
 - Two major implementations: apparmor or selinux

Implementations

Overview over implementations (1/2)

- “System” container engines (runs multiple services like a VM)
 - Systemd-nspawn, lxc/lxd, Openvz, ...



LXD



OpenVZ




- Application sandboxes, special purpose container engines
 - Bubblewrap, snap, chromium sandbox
 - Singularity (High performance Computing)




Overview over implementations (2/2)

- “Application” container engines (one container per service)
 - **Open Container Initiative:** Mostly standardized images/runtime
 - Docker, Podman, Containerd, ...



- Cluster manager:
 - Schedule container over multiple hosts
 - make use of the underlying container engine



Nomad



kubernetes



Used to run your tests!

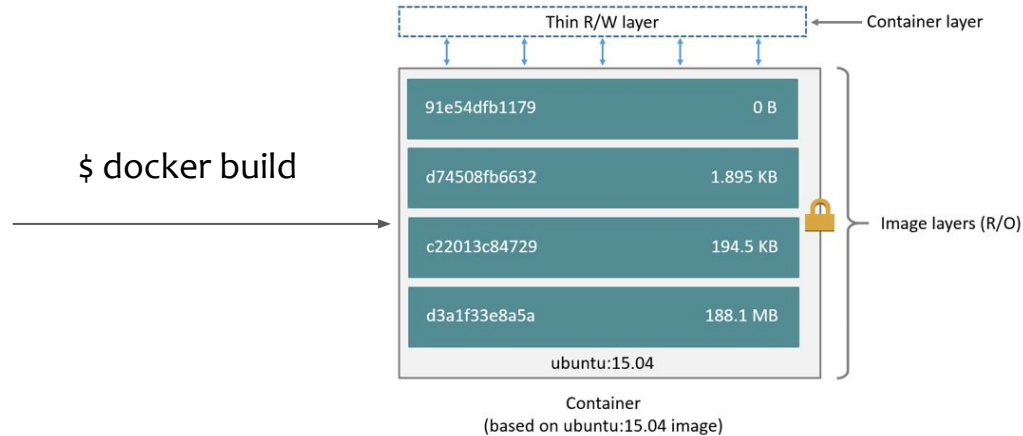
OCI image & runtime

- **Docker:** Most popular container runtime
- **Dockerfile:** build description to build container images
- Makes heavy of overlay/snapshot filesystems to save space

Dockerfile*

```
FROM ubuntu:18.04  
COPY . /app  
RUN make /app  
CMD python /app/app.py
```

\$ docker build



- Standardized by OCI:
 - both runtime (i.e VMs instead of container) and build tools (buildah instead of Dockerfile) can be exchanged

* Dockerfile used in our course <https://gist.github.com/Mic92/7ae91f5f0239acb56c67535c683f1bc1>

Demo time!